

Problem 2: Neural Networks (50 points)

Part A: Trial by Big Head (15 points)

You've just been hired by tech-giant Hooli to lead the AI lab. A fellow team member affectionately known as Big Head decides to put your skills to the test with a series of conceptual questions. For each of the following questions, circle the **one** best answer:

A1 (3 points) Your neural net is still inaccurate after having been trained for a fixed number of iterations. You consider adjusting the learning rate. Which of the following adjustments might cause the network to converge to a more accurate solution?

- A) Increasing the learning rate.
- B) Decreasing the learning rate.
- C) Either increasing or decreasing the learning rate.
- D) None of these: the change in accuracy is independent of learning rate.

A2 (3 points) One motivation for using the sigmoid instead of the stairstep threshold function in back propagation is:

- A) It's computationally inefficient to compute the stairstep gradient.
- B) The stairstep function is an odd function.
- C) Back propagation is only well-defined for the sigmoid function.
- D) The sigmoid function is differentiable everywhere.
- E) None of the above.


A3 (3 points) Suppose you decide to use **the identity function** as your threshold function:

$$f(x) = x$$

In other words, the output of each neuron is just the weighted sum of its inputs. Of the following three neural networks, which one can accurately classify new inputs **that the others cannot**?

- A) *Network A*: A neural net with 2 input neurons and 1 output neuron.
- B) *Network B*: A neural net with 2 input neurons, 1 hidden-layer neuron, and 1 output neuron.
- C) *Network C*: A neural net with 2 input neurons, 3 hidden-layer neurons, and 1 output neuron.
- D) Networks A, B, and C can accurately classify the same inputs.
- E) Can't tell without more information.

A4 (6 points) You're building a neural net to classify the positive and negative data samples shown below on the left: positive samples will output 1, and negative samples will output 0. Each sample has features x and y , however, you are only allowed to use the feature $p = x \cdot y$ as input to any neural network. (See below on the right for an illustration of p 's value as a function of x and y .)



For each of the following network architectures, could a neural network of that shape correctly classify the data? (Assume all neurons use the stairstep threshold function.) In each row, circle either **YES** or **NO**:

Network 1: $p \rightarrow \left(\int \right) \rightarrow \text{out}$ YES NO


Network 2: $p \rightarrow \left(\int \right) \rightarrow \left(\int \right) \rightarrow \text{out}$ YES NO

Network 3: $p \rightarrow \left(\int \right) \rightarrow \left(\int \right) \rightarrow \text{out}$ YES NO

Part B: Let's Go! (21 points)

Impressed with your knowledge, Big Head places you on a top-secret team tasked with building a Go-playing bot called BetaGo. Your job is to build a neural net that can be used to classify winning (\oplus) and losing (\ominus) board positions, as shown on the right.

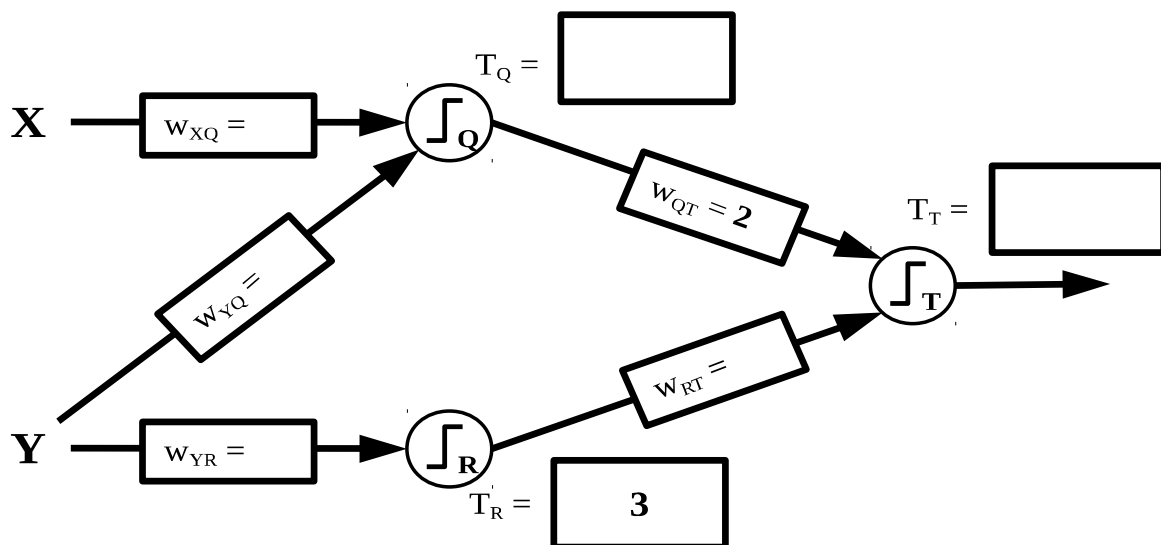
Due to budget cuts, you are limited to using a small net with only three neurons (each of which uses a stairstep threshold function).



Big Head tells you that the small network won't be able to classify *all* of the samples correctly: in fact, it will misclassify *one* sample.

In the neural net skeleton below, we have already filled in weight w_{QT} and threshold T_R for you. Fill in the remaining **four** weights and **two** thresholds with INTEGER values such that the net will correctly classify all but one sample. (We have provided space on the next page to show your work for partial credit.)

Note: A Win outputs 1; a Loss outputs 0.



For partial credit for part B, you can show your work here:


Part C: Back-handed Back-prop (14 points)

Downtrodden startup founder Richard Hendrix approaches and tells Big Head that he doesn't need to manually pick all the weights: there's an automatic way to train the weights of the network using back propagation! Big Head asks for a demonstration of back propagation.

For your convenience, an equation sheet for neural nets is provided on a tear-off sheet at the end of the quiz.

C1 (8 points) The following neural network uses the **sigmoid threshold function**. The weights have been initialized randomly, and you have just performed forward propagation to determine the current outputs of each of the network's neurons. Calculate the values of δ_T and δ_Q , assuming the following:

$$X = 3, \quad Y = 7, \quad \text{Learning Rate} = 1, \quad \text{Desired Output} = 0$$



Space is provided on the next page to show your work.

$\delta_T =$

$\delta_Q =$

For partial credit for part C1, you can show your work here:

C2 (6 points) Calculate the weight update for W_{YR}^{new} , the weight between input Y and neuron R. **You can leave your answer in terms of δ_Q , δ_R , and δ_T .**

$W_{YR}^{new} =$

For partial credit for part C2, you can show your work here: